

Loading the Dataset

Finding the best agent for a company using the past recruitment data

```
In [1]: import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [2]: df = pd.read_csv('C:\\Users\\SIGMA\\Desktop\\Machine Learning\\10. Logistic Regression\\Module 10\\helper text\\3. Implemen
df.shape
```

Out[2]: (8844, 19)

```
In [3]: df.head()
```

Out[3]:

| | ID | Office_PIN | Applicant_City_PIN | Applicant_Gender | Applicant_Marital_Status | Applicant_Occupation | Applicant_Qualification | Manager_Jr |
|---|------------|------------|--------------------|------------------|--------------------------|----------------------|-------------------------|------------|
| 0 | FIN1000001 | 842001 | 844120 | M | M | Others | Graduate | |
| 1 | FIN1000002 | 842001 | 844111 | M | S | Others | Class XII | |
| 2 | FIN1000003 | 800001 | 844101 | M | M | Business | Class XII | |
| 3 | FIN1000004 | 814112 | 814112 | M | S | Salaried | Class XII | |
| 4 | FIN1000005 | 814112 | 815351 | M | M | Others | Class XII | |

```
In [4]: df.columns
```

```
Out[4]: Index(['ID', 'Office_PIN', 'Applicant_City_PIN', 'Applicant_Gender',  
             'Applicant_Marital_Status', 'Applicant_Occupation',  
             'Applicant_Qualification', 'Manager_Joining_Designation',  
             'Manager_Current_Designation', 'Manager_Grade', 'Manager_Status',  
             'Manager_Gender', 'Manager_Num_Application', 'Manager_Num_Coded',  
             'Manager_Business', 'Manager_Num_Products', 'Manager_Business2',  
             'Manager_Num_Products2', 'Business_Sourced'],  
            dtype='object')
```

Imputing Missing Values

```
In [5]: df.isnull().sum()
```

```
Out[5]: ID                                0  
Office_PIN                               0  
Applicant_City_PIN                       0  
Applicant_Gender                         53  
Applicant_Marital_Status                 59  
Applicant_Occupation                    1090  
Applicant_Qualification                   71  
Manager_Joining_Designation               0  
Manager_Current_Designation              0  
Manager_Grade                           0  
Manager_Status                           0  
Manager_Gender                           0  
Manager_Num_Application                   0  
Manager_Num_Coded                        0  
Manager_Business                         0  
Manager_Num_Products                     0  
Manager_Business2                        0  
Manager_Num_Products2                    0  
Business_Sourced                         0  
dtype: int64
```

```
In [7]: df[['Applicant_Gender', 'Applicant_Marital_Status', 'Applicant_Occupation', 'Applicant_Qualification']].head()
```

```
Out[7]:
```

| | Applicant_Gender | Applicant_Marital_Status | Applicant_Occupation | Applicant_Qualification |
|---|------------------|--------------------------|----------------------|-------------------------|
| 0 | M | M | Others | Graduate |
| 1 | M | S | Others | Class XII |
| 2 | M | M | Business | Class XII |
| 3 | M | S | Salaried | Class XII |
| 4 | M | M | Others | Class XII |

1. Missing Values in Applicant Gender

```
In [8]: #checking value_Counts  
df['Applicant_Gender'].value_counts()
```

```
Out[8]: M    6656  
       F    2135  
       Name: Applicant_Gender, dtype: int64
```

```
In [9]: #imputing missing with mode  
df['Applicant_Gender'].fillna('M', inplace=True)
```

2. Missing Values in Applicant Marital Status

```
In [10]: #checking value_Counts  
df['Applicant_Marital_Status'].value_counts()
```

```
Out[10]: M    5733  
        S    3042  
        W      6  
        D      4  
        Name: Applicant_Marital_Status, dtype: int64
```

```
In [11]: #imputing missing with mode
df['Applicant_Marital_Status'].fillna('M', inplace=True)
```

3. Missing Values in Applicant Occupation

```
In [12]: #checking value_Counts
df['Applicant_Occupation'].value_counts()
```

```
Out[12]: Salaried          3546
Business       2157
Others         1809
Self Employed   146
Student         96
Name: Applicant_Occupation, dtype: int64
```

```
In [13]: #imputing missing with mode
df['Applicant_Occupation'].fillna('Salaried', inplace=True)
```

4. Missing Values in Applicant Qualification

```
In [14]: #checking value_Counts
df['Applicant_Qualification'].value_counts()
```

```
Out[14]: Class XII          5426
Graduate          2958
Class X           195
Others            116
Masters of Business Administration    71
Associate / Fellow of Institute of Chartered Accountants of India    3
Professional Qualification in Marketing    1
Associate/Fellow of Insurance Institute of India    1
Associate/Fellow of Acturial Society of India    1
Associate/Fellow of Institute of Company Secretories of India    1
Name: Applicant_Qualification, dtype: int64
```

```
In [15]: #imputing missing with mode  
  
df['Applicant_Qualification'].fillna('Class XII', inplace=True)
```

```
In [16]: df.isnull().sum()
```

```
Out[16]: ID                                0  
Office_PIN                                0  
Applicant_City_PIN                        0  
Applicant_Gender                         0  
Applicant_Marital_Status                 0  
Applicant_Occupation                     0  
Applicant_Qualification                  0  
Manager_Joining_Designation              0  
Manager_Current_Designation              0  
Manager_Grade                            0  
Manager_Status                           0  
Manager_Gender                           0  
Manager_Num_Application                  0  
Manager_Num_Coded                        0  
Manager_Business                         0  
Manager_Num_Products                     0  
Manager_Business2                        0  
Manager_Num_Products2                    0  
Business_Sourced                         0  
dtype: int64
```

Dealing with Categorical Variables

```
In [17]: df.dtypes
```

```
Out[17]: ID                                object
Office_PIN                               int64
Applicant_City_PIN                       int64
Applicant_Gender                         object
Applicant_Marital_Status                 object
Applicant_Occupation                     object
Applicant_Qualification                   object
Manager_Joining_Designation              object
Manager_Current_Designation              object
Manager_Grade                           float64
Manager_Status                           object
Manager_Gender                           object
Manager_Num_Application                  float64
Manager_Num_Coded                       float64
Manager_Business                         float64
Manager_Num_Products                     float64
Manager_Business2                       float64
Manager_Num_Products2                   float64
Business_Sourced                         int64
dtype: object
```

```
In [38]: df.columns
```

```
Out[38]: Index(['ID', 'Office_PIN', 'Applicant_City_PIN', 'Applicant_Gender',
               'Applicant_Marital_Status', 'Applicant_Occupation',
               'Applicant_Qualification', 'Manager_Joining_Designation',
               'Manager_Current_Designation', 'Manager_Grade', 'Manager_Status',
               'Manager_Gender', 'Manager_Num_Application', 'Manager_Num_Coded',
               'Manager_Business', 'Manager_Num_Products', 'Manager_Business2',
               'Manager_Num_Products2', 'Business_Sourced'],
              dtype='object')
```

```
In [18]: categorical_cols = ['Applicant_Gender', 'Applicant_Marital_Status', 'Applicant_Occupation', 'Applicant_Qualification',
                             'Manager_Joining_Designation', 'Manager_Current_Designation', 'Manager_Status', 'Manager_Gender']
```

```
for i in categorical_cols:
    print('*****', i, '*****')
    print(df[i].value_counts())
    print('')
```

```
***** Applicant_Gender *****
```

```
M      6709
```

```
F      2135
```

```
Name: Applicant_Gender, dtype: int64
```

```
***** Applicant_Marital_Status *****
```

```
M      5792
```

```
S      3042
```

```
W         6
```

```
D         4
```

```
Name: Applicant_Marital_Status, dtype: int64
```

```
***** Applicant_Occupation *****
```

```
Salaried      4636
```

```
Business      2157
```

```
Others        1809
```

```
Self Employed   146
```

```
Student         96
```

```
Name: Applicant_Occupation, dtype: int64
```

```
***** Applicant_Qualification *****
```

```
Class XII                                     5497
```

```
Graduate                                    2958
```

```
Class X                                     195
```

```
Others                                      116
```

```
Masters of Business Administration          71
```

```
Associate / Fellow of Institute of Chartered Accountants of India    3
```

```
Professional Qualification in Marketing      1
```

```
Associate/Fellow of Insurance Institute of India                      1
```

```
Associate/Fellow of Actuarial Society of India                        1
```

```
Associate/Fellow of Institute of Company Secretaries of India        1
```

```
Name: Applicant_Qualification, dtype: int64
```

```

***** Manager_Joining_Designation *****
Level 1      4632
Level 2      2787
Level 3      1146
Level 4       200
Other         58
Level 6       18
Level 7        2
Level 5        1
Name: Manager_Joining_Designation, dtype: int64

***** Manager_Current_Designation *****
Level 2      3208
Level 1      2479
Level 3      2033
Level 4      1031
Level 5        93
Name: Manager_Current_Designation, dtype: int64

***** Manager_Status *****
Confirmation   5277
Probation     3567
Name: Manager_Status, dtype: int64

***** Manager_Gender *****
M      7627
F      1217
Name: Manager_Gender, dtype: int64

```

```
In [19]: df = pd.get_dummies(df)
```

Logistic regression

Train Test split


```
In [23]: x = df.drop(['Business_Sourced'],axis=1)
y = df['Business_Sourced']
```

```
In [51]: from sklearn.model_selection import train_test_split
train_x, valid_x, train_y, valid_y= train_test_split(x, y, test_size = 0.3, random_state=1)
```

Linear regression Model

```
In [113]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import f1_score, auc
from sklearn.metrics import roc_auc_score
```

```
In [53]: logreg = LogisticRegression()
logreg.fit(train_x, train_y)
```

/home/aishwarya/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:432: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.
FutureWarning)

```
Out[53]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, l1_ratio=None, max_iter=100,
multi_class='warn', n_jobs=None, penalty='l2',
random_state=None, solver='warn', tol=0.0001, verbose=0,
warm_start=False)
```

```
In [115]: pred_train = logreg.predict_proba(train_x)
pred_valid = logreg.predict_proba(valid_x)
```

```
In [117]: roc_auc_score(train_y, pred_train[:,1])
```

```
Out[117]: 0.4723336326910138
```

```
In [114]: roc_auc_score(valid_y, pred_valid[:,1])
```

```
Out[114]: 0.4697613206972208
```

